

Design and Implementation of a Multipath Extension for QUIC



TECHNISCHE
UNIVERSITÄT
DARMSTADT


Tobias Viernickel
Betreuer: Alexander Frömmgen

Prof. Dr.-Ing. Ralf Steinmetz
KOM - Multimedia Communications Lab

Motivation

Single-Path protocols


- TCP
- UDP
- QUIC



Motivation

Multipath protocols

- MPTCP
- MPUDP
- SCTP



The QUIC Transport Protocol: Design and Internet-Scale Deployment

Adam Langley, Alistair Riddoch, Alyssa Wilk, Antonio Vicente, Charles Krasic, Dan Zhang, Fan Bailey, Jeremy Dorfman, Jim Roskind, Rade, Ryan Hamilton, Victor Vasiliev, Shi *

*“... low-latency transport protocol ...
enable rapid deployment ...
continued evolution of transport mechanisms ...
deployed at Google on thousands of servers ...”^[1]*

No Multipath Support

Describe our motivation for developing a new transport, the principles that guided our design, the Internet-scale process that we used to perform iterative experiments on QUIC, performance improvements seen by our various services, and our experience deploying QUIC globally. We also share lessons about transport design and the Internet ecosystem that we learned from our deployment.

CCS CONCEPTS

• Networks → Network protocol design; Transport protocols; Cross-layer protocols;

ACM Reference format:

Adam Langley, Alistair Riddoch, Alyssa Wilk, Antonio Vicente, Charles Krasic, Dan Zhang, Fan Yang, Fedor Kouranov, Ian Swett, Janardhan Iyengar, Jeff Bailey, Jeremy Dorfman, Jim Roskind, Joanna Kulik, Patrik Westin, Raman Tenneneti, Robbie Shade, Ryan Hamilton, Victor Vasiliev, Wan-Teh




Figure 1: QUIC in the traditional HTTPS stack.

TCP (Figure 1). We developed QUIC as a user-space transport with UDP as a substrate. Building QUIC in user-space facilitated its deployment as part of various applications and enabled iterative changes to occur at application update timescales. The use of UDP allows QUIC packets to traverse middleboxes. QUIC is an encrypted transport: packets are authenticated and encrypted, preventing modification and limiting ossification of the protocol by middleboxes. QUIC uses a cryptographic handshake that minimizes handshake latency for most connections by using known server credentials on repeat connections and by removing redundant handshake-overhead at multiple layers in the network stack. QUIC eliminates head-of-line blocking delays by using a lightweight data-structuring abstraction, streams, which are multiplexed within a single connection so that

Can we make QUIC multipath capable?

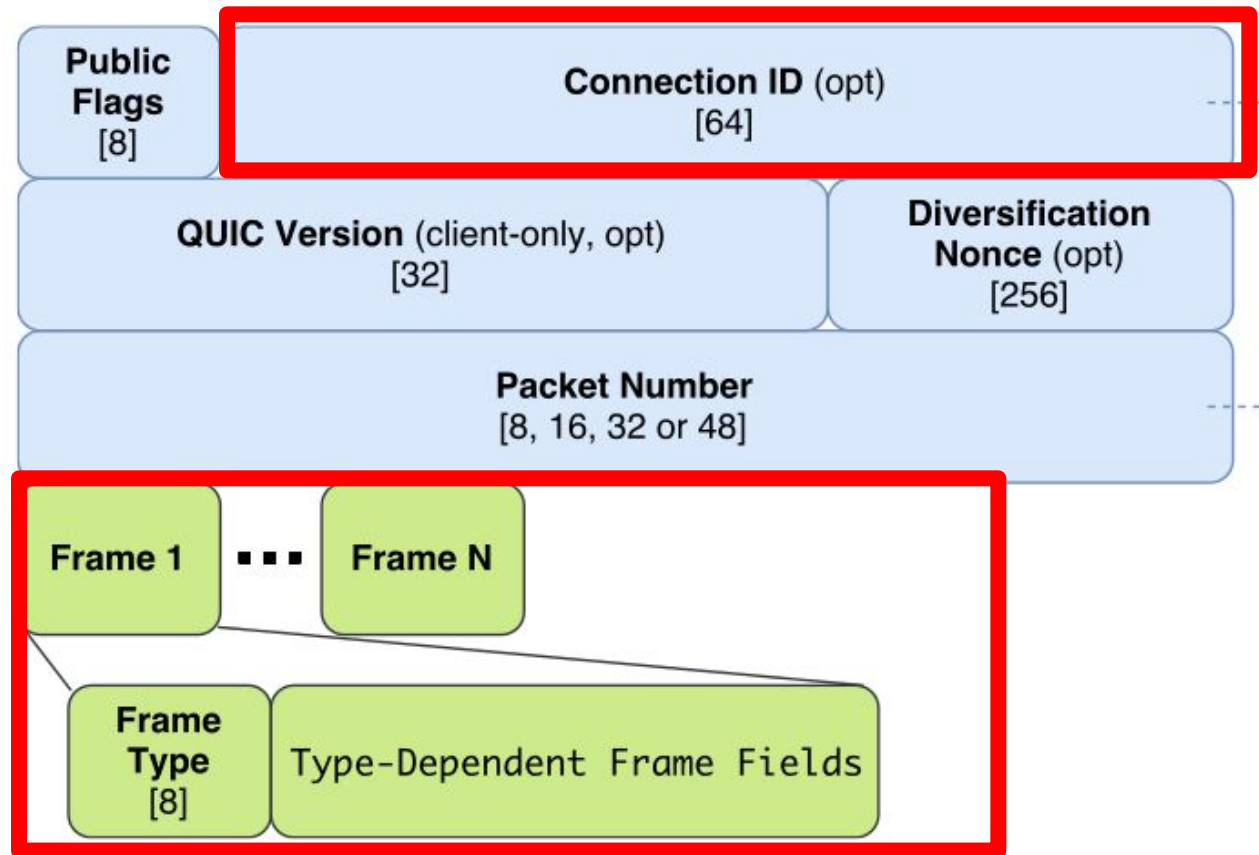
Can MPQUIC improve the performance?

Has MPQUIC any fundamental advantages?

Background: QUIC




TECHNISCHE
UNIVERSITÄT
DARMSTADT




The Design of MPQUIC

Network Stack Comparison


a) traditional QUIC



b) MPTCP

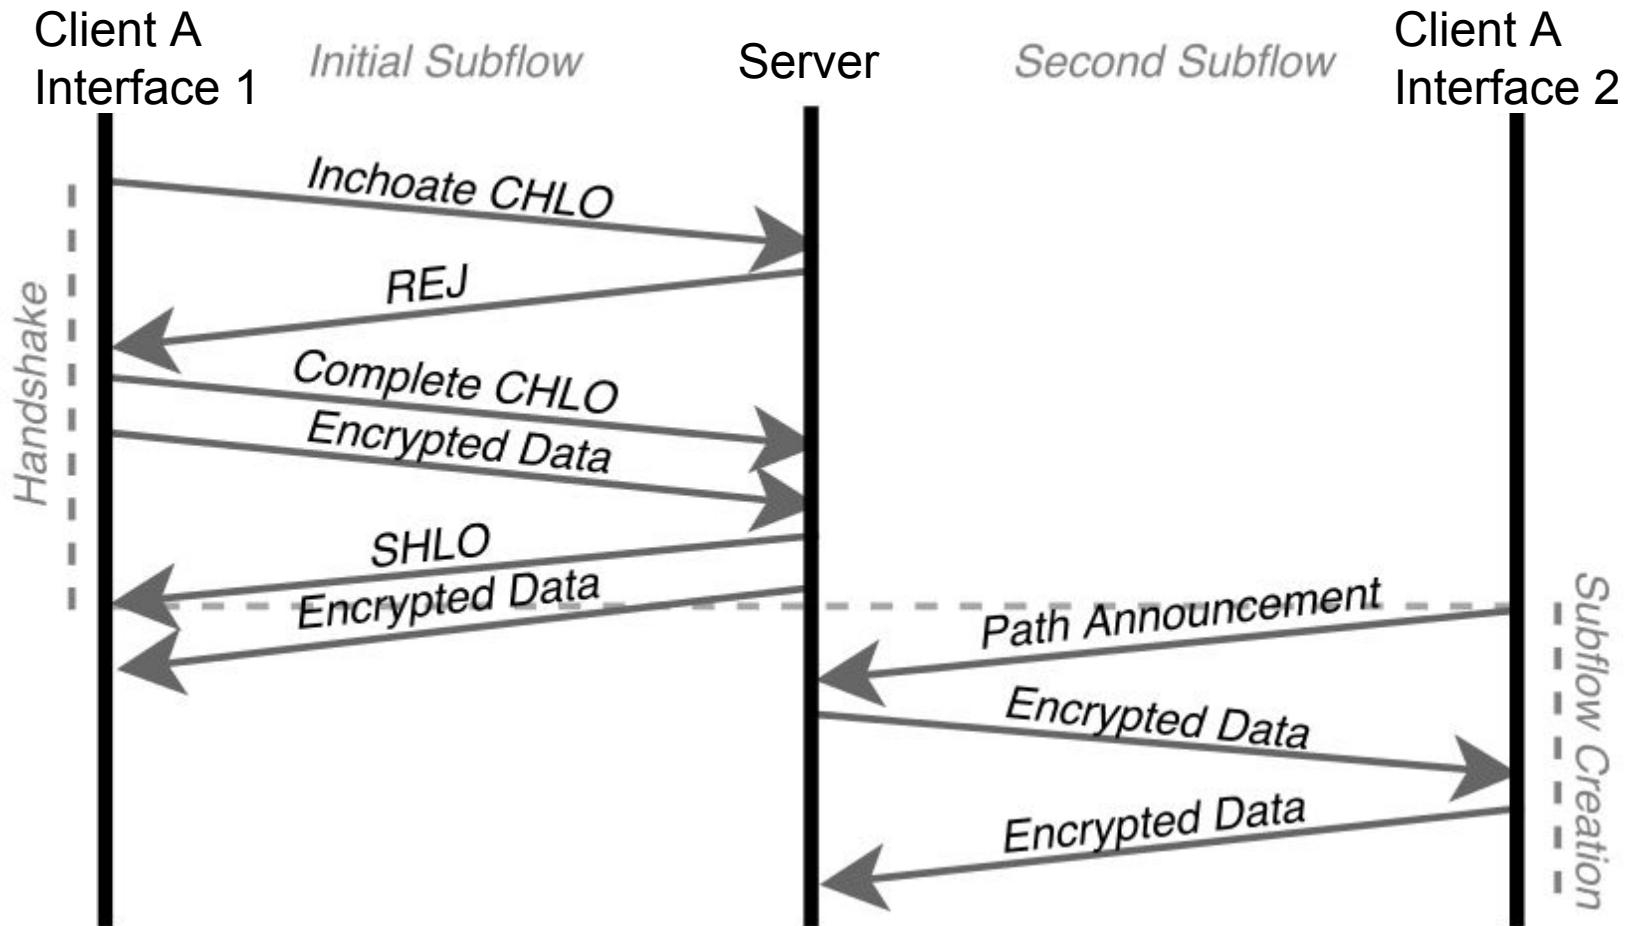


c) MPQUIC



The Design of MPQUIC

Subflow Establishment: Announcement

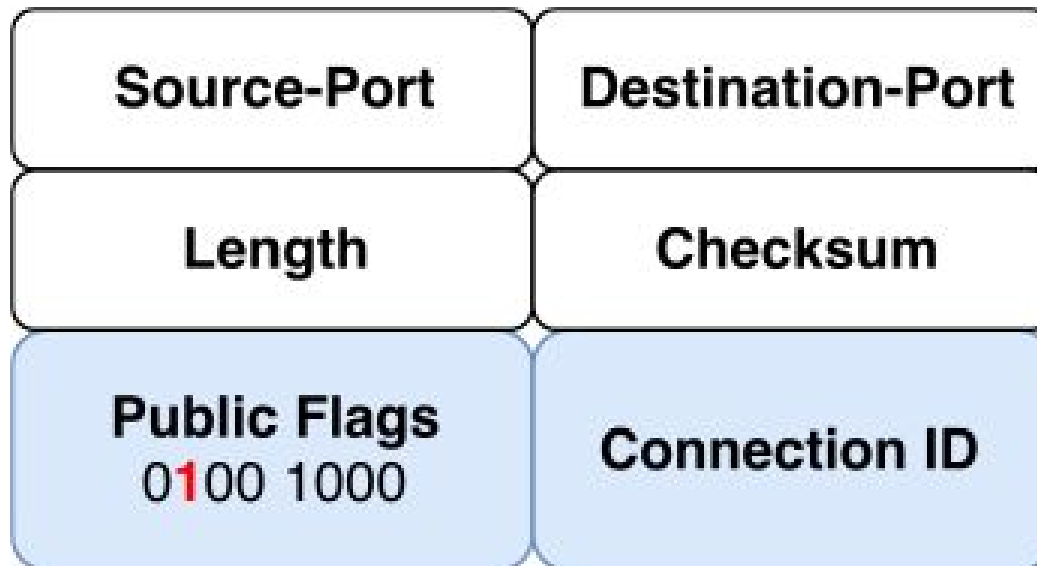


The Design of MPQUIC

Subflow Establishment: Announcement

The Announcement Packet

- IP/Port inside the UDP/IP Header
- ConnectionID
- MultipathFlag

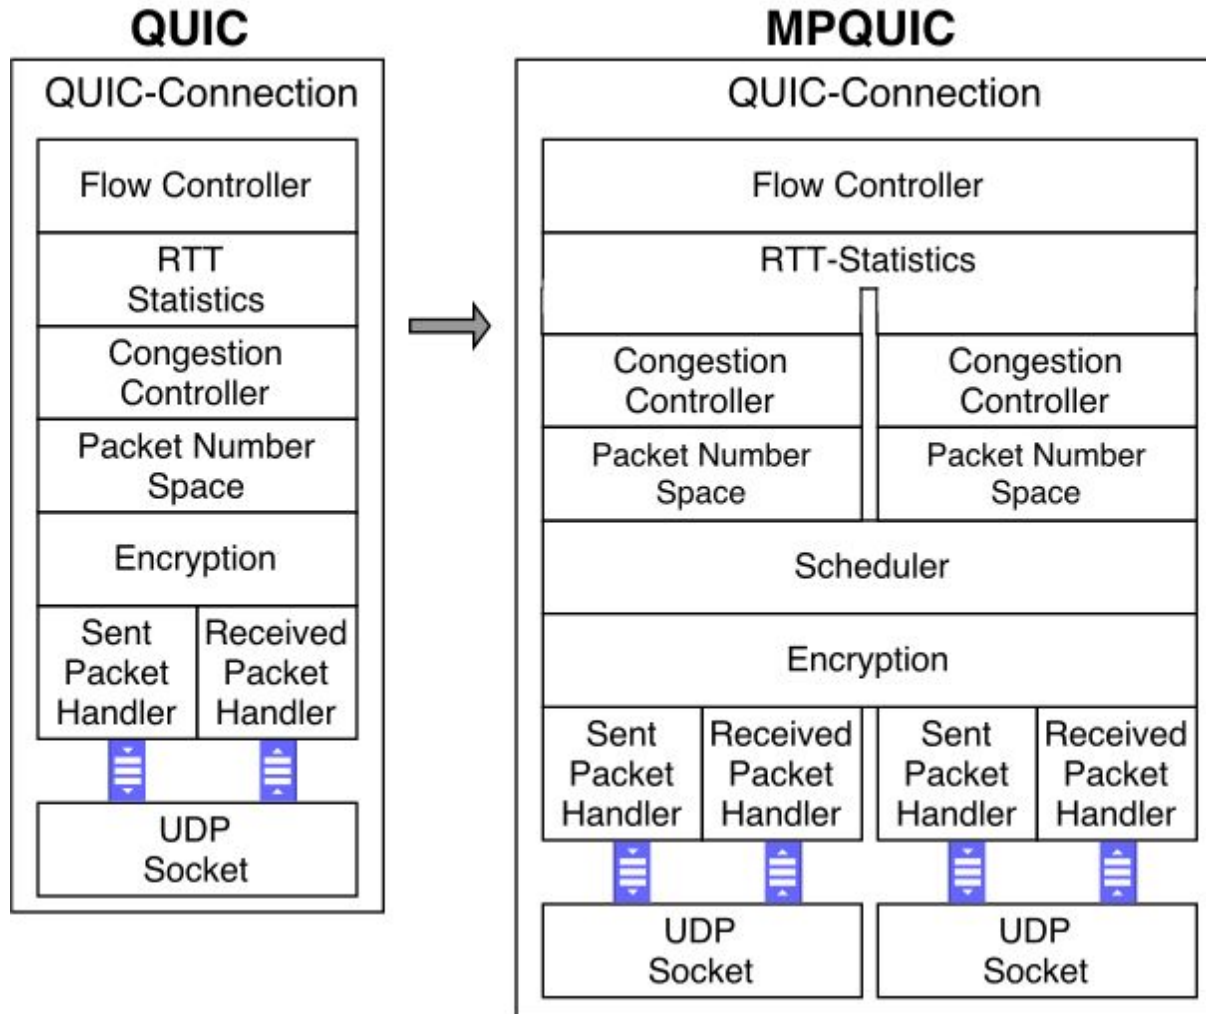


The Design of MPQUIC

Specifying the individual Components




TECHNISCHE
UNIVERSITÄT
DARMSTADT




Evaluation

General Setup using Mininet

Single-Path Setup



Multipath Setup



Evaluation

Speedup of multipath QUIC

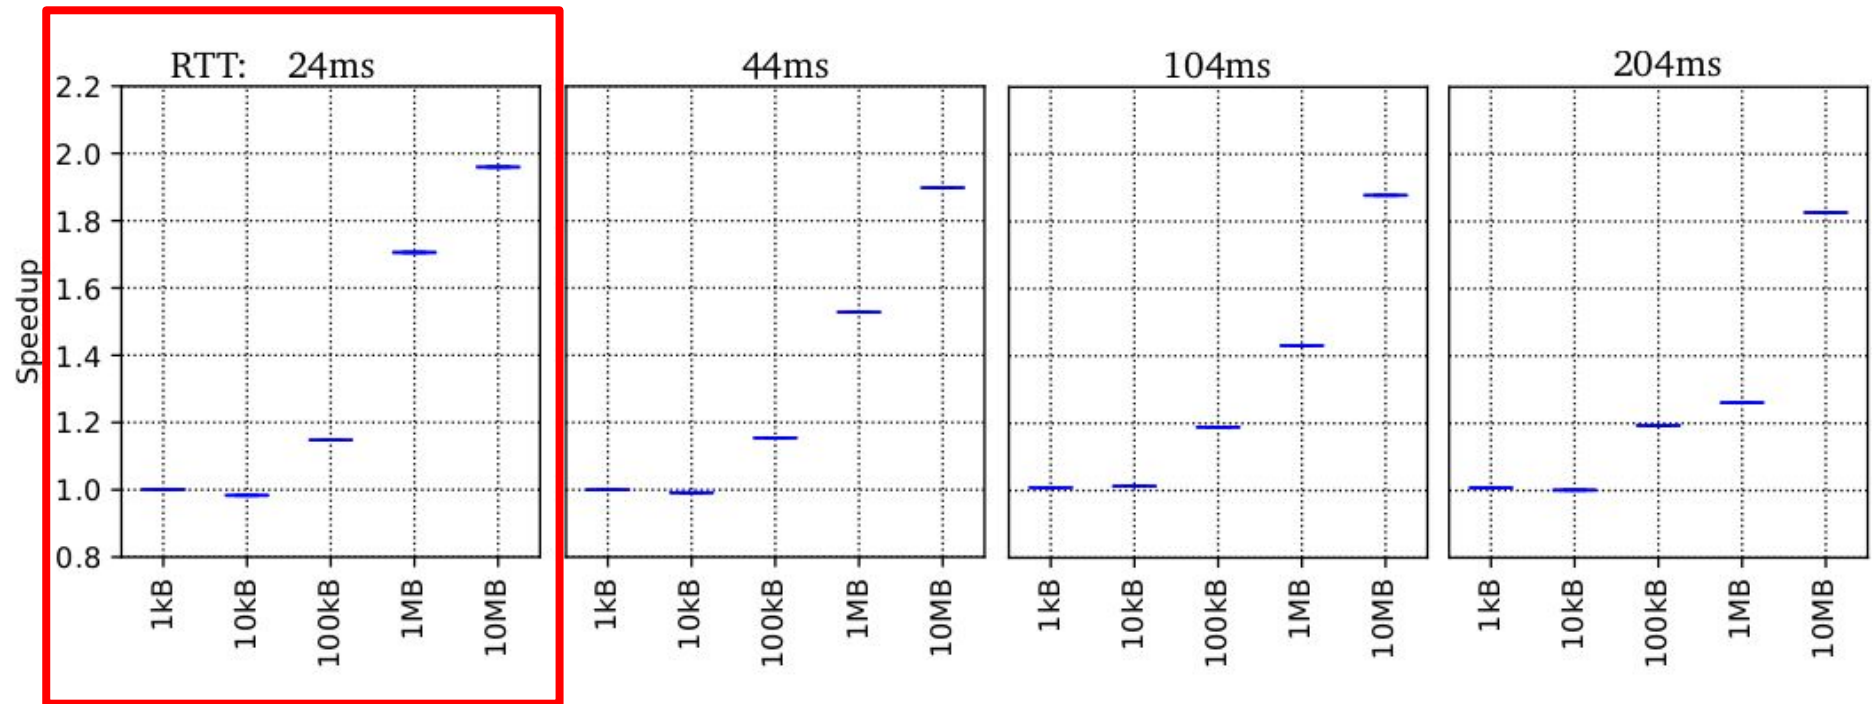
HTTP/2 file requests of different sizes

Initial Path:

Bandwidth: 10Mbps

Second Path:

Bandwidth: 10Mbps



Evaluation

Speedup of multipath QUIC

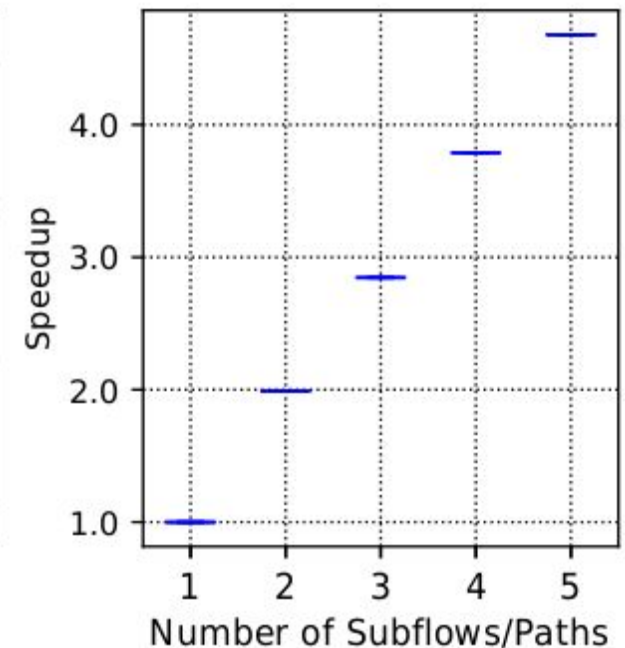
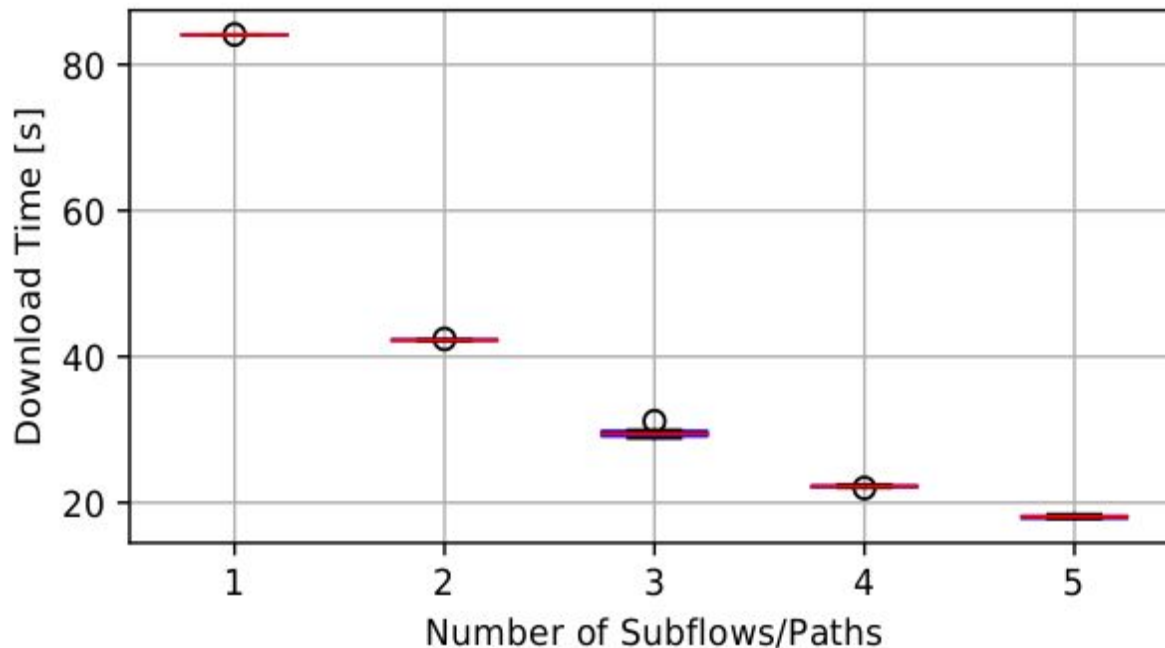
HTTP/2 file requests for a various number of paths

- File size: 100 MB

All Paths:

Bandwidth: 10Mbps

RTT: 44ms



Evaluation

Comparison of MPQUIC, QUIC, MPTCP, TCP

HTTP/2 file request

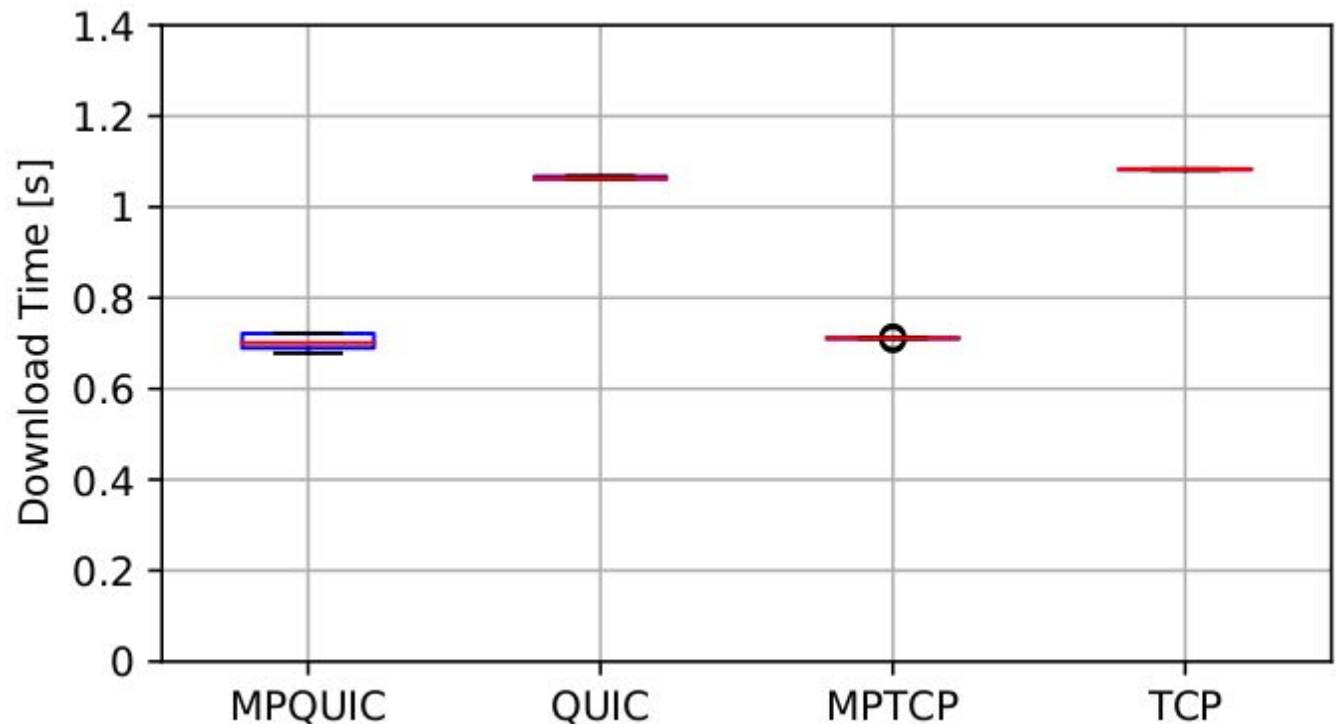
- File size: 1 MB

Initial Path

Bandwidth: 10Mbps
RTT: 44ms

Second Path

Bandwidth: 10Mbps
RTT: 44ms



Evaluation

Priority Scheduler

HTTP/2 file requests


- Low Priority File: 10 MB
- High Priority File: 1 KB (request after ~500ms)

Initial Path

Bandwidth: 10Mbps
RTT: 44ms

Second Path

Bandwidth: 10Mbps
RTT: 104ms



Can we make QUIC multipath capable?



Can MPQUIC improve the performance?



Has MPQUIC any fundamental advantages?



- *Deployability*
- *Middlebox interference*
- *Stream to Subflow Scheduling*

Outlook

- **Complete the design of MPQUIC**
- **Wide deployment and evaluation of various application scenarios**
- **Optimized stream to subflow scheduling (e.g., priorities)**

Thank you for your attention! Questions?





TECHNISCHE
UNIVERSITÄT
DARMSTADT




Evaluation

Experimental Priority Scheduler

MPQUIC



MPTCP




Evaluation

Comparison of MPQUIC, QUIC, MPTCP, TCP

HTTP/2 file requests of small sized files

- Bandwidth(const): 10 Mbps
- RTT(const):44ms

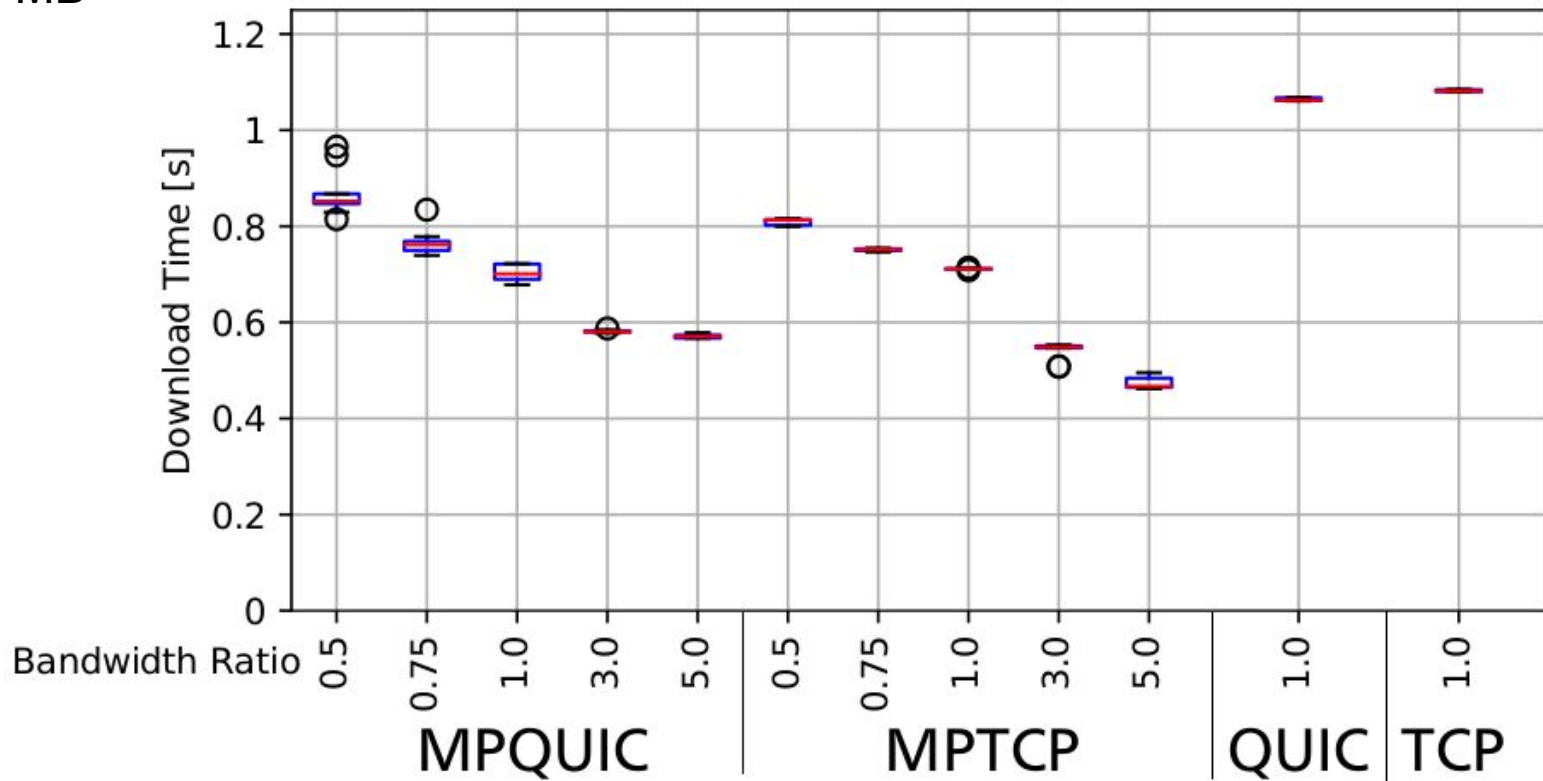


Evaluation

Comparison of MPQUIC, QUIC, MPTCP, TCP

HTTP/2 file requests for various bandwidth ratios

- Initial path bandwidth: 10 Mbps
- RTT(const): 44ms
- 1 MB

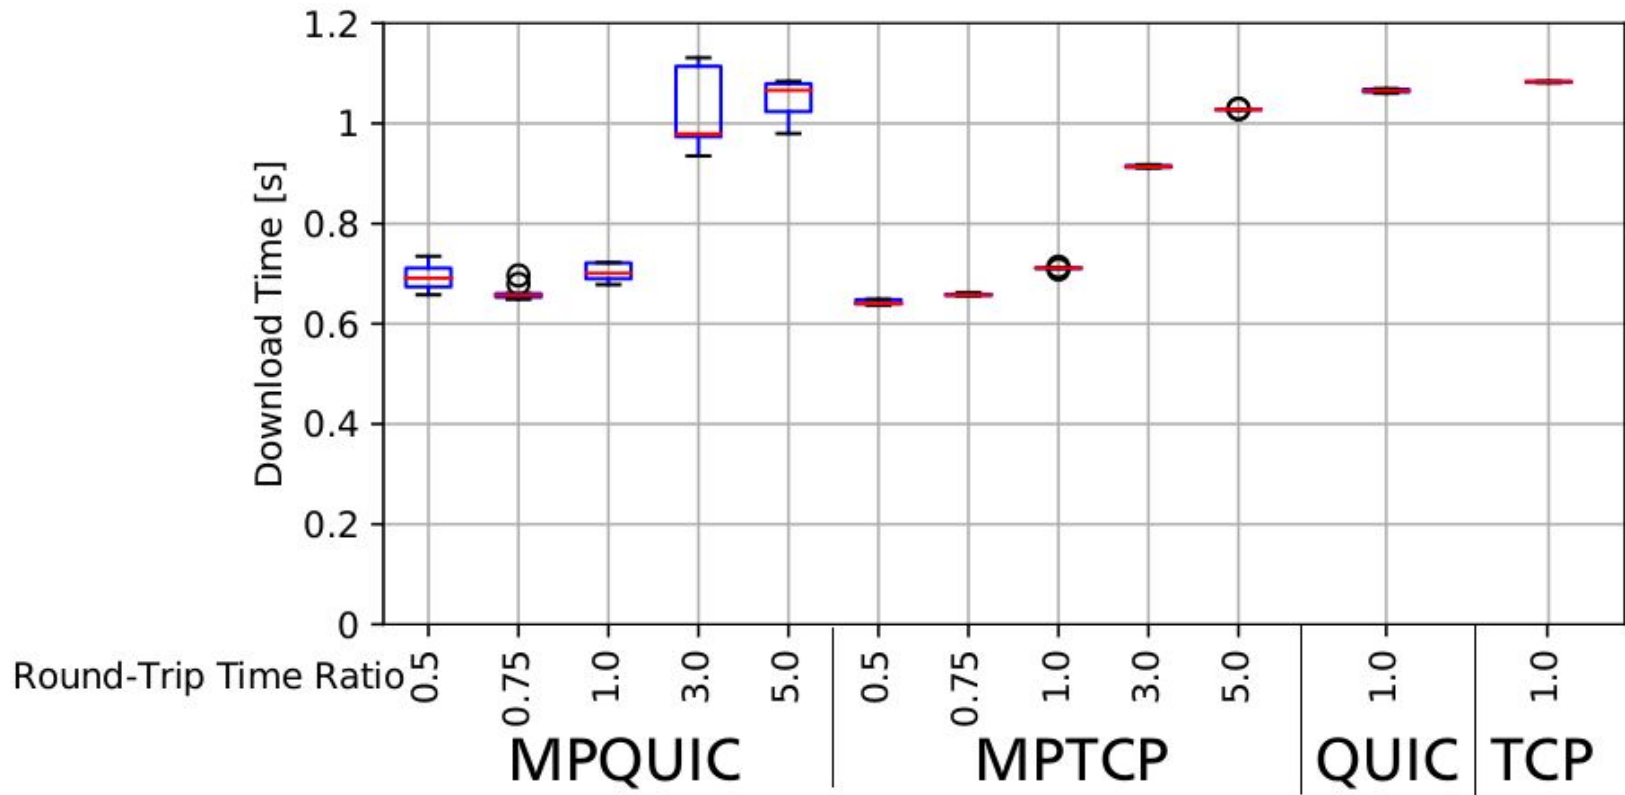


Evaluation

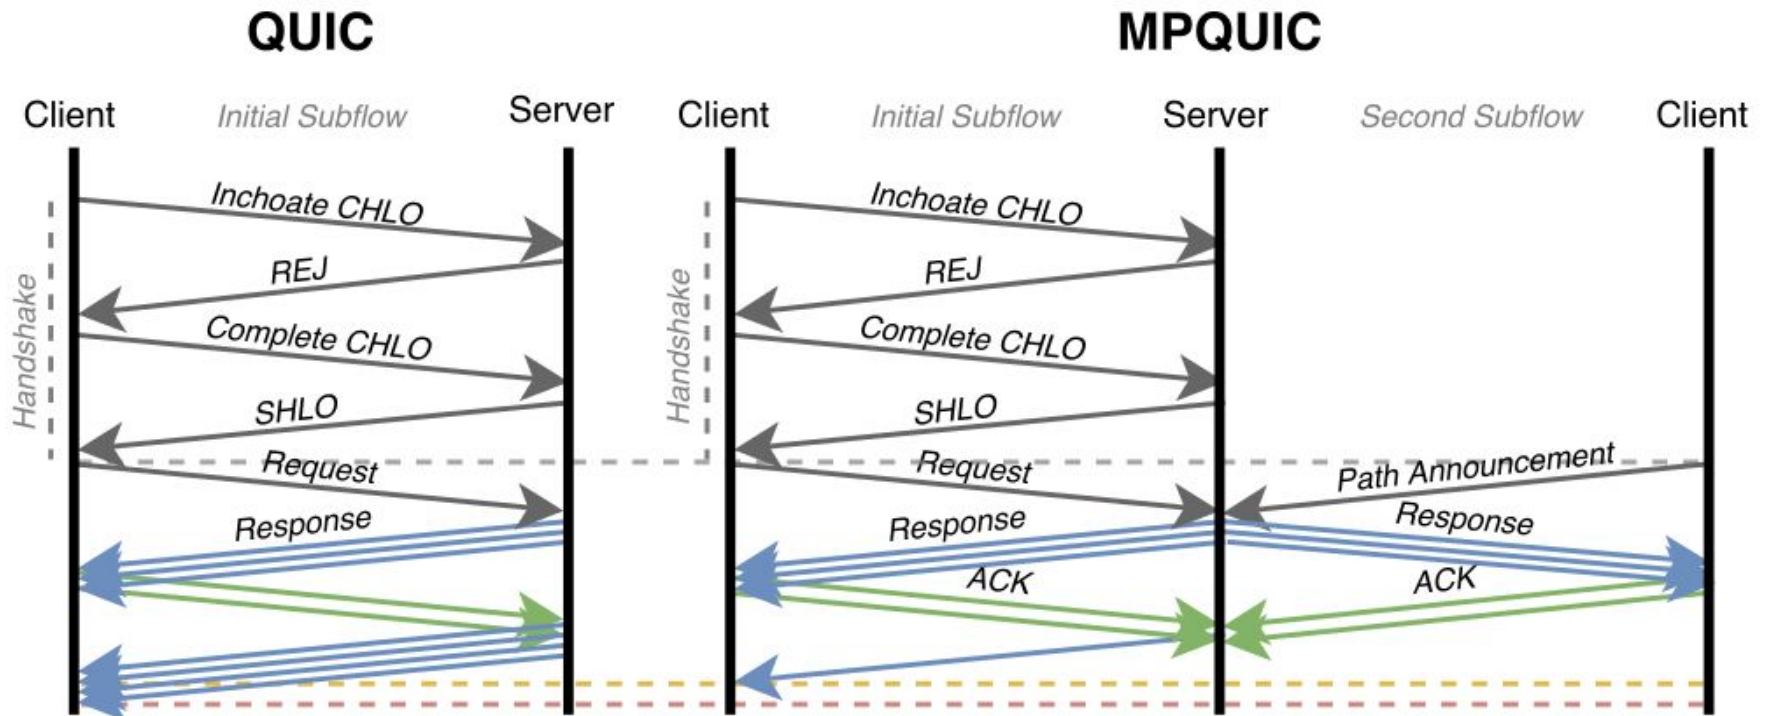
Comparison of MPQUIC, QUIC, MPTCP, TCP

HTTP/2 file requests for various RTT ratios

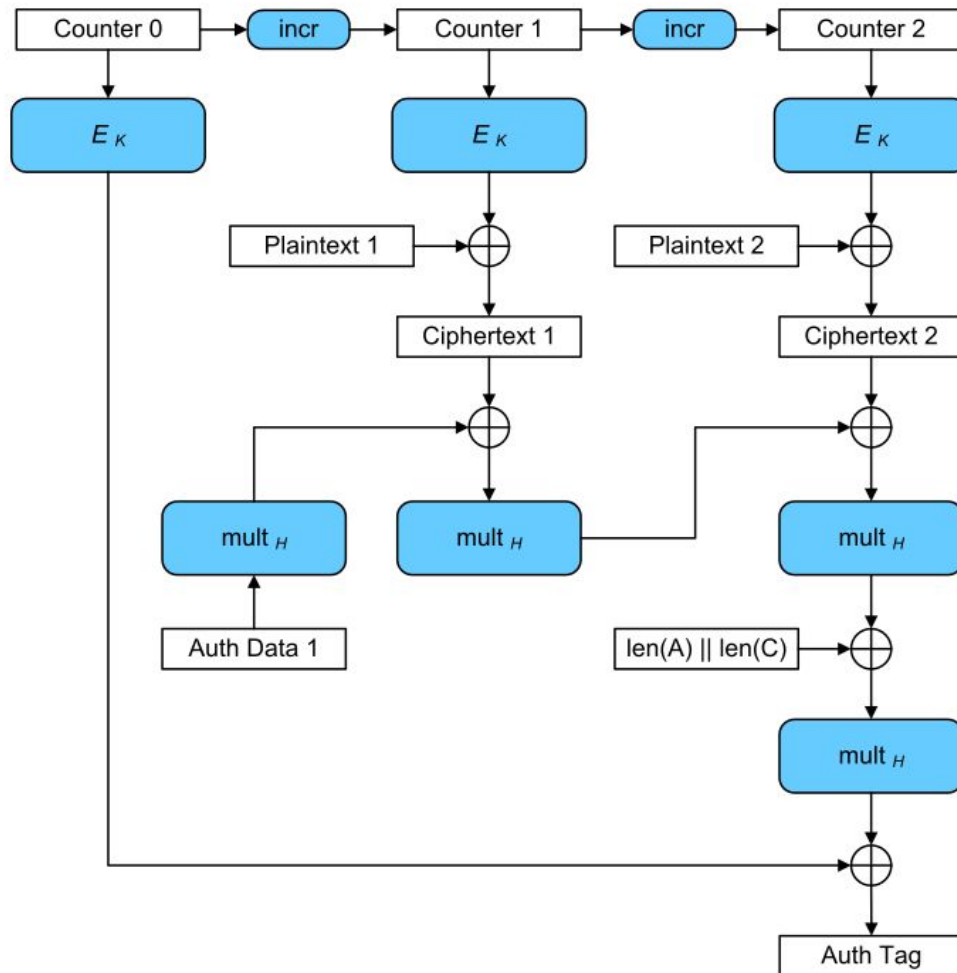
- Bandwidth(const): 10 Mbps
- Initial path RTT: 44ms
- 1 MB



30 KB file download



Encryption



Public Flags

Public Flag	Meaning
0x01	Public Flag Version
0x02	Public Reset Flag
0x04	Presence of Diversification Nonce
0x08	Presence of 8 byte Connection ID
0x30	(2 bit mask) Indicates length of Packet Number
0x40	Reserved for multipath use
0x80	Not used



$$RTT_1 = (r_1 - (r_0 - RTT_0 \times \frac{1}{2})) \times 2$$

Implementation of MPQUIC in quic-go

Extension for quic-go

Adjustments for Evaluation

- Initial congestion window size (10 packets)
- Default accept source-address-token
- Simple HTML Web-Scraper
- HTTP/2 TCP/TLS server and client

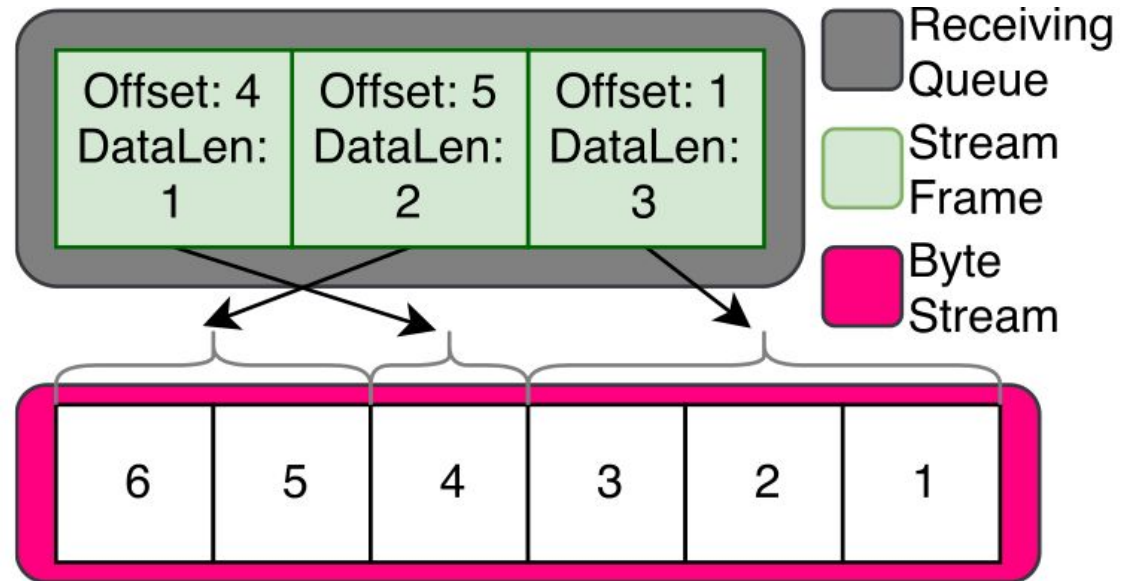
Plug and Play Scheduler

- Declaration of a stream and subflow scheduler lambda function
- Enabled easy modification of the scheduler
- Select scheduler via command-line parameter

Background: QUIC Streams

QUIC Streams:

- Independent sequences of data
- Multiple Streams per Connection
 - Identifier: StreamID
 - In-order delivery
 - Offset



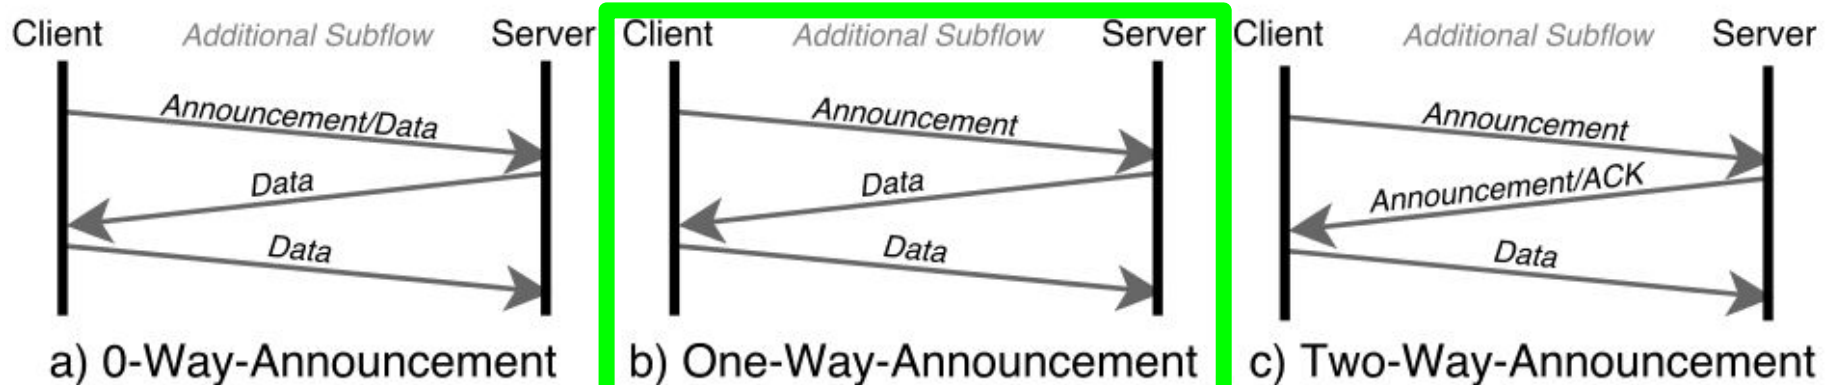
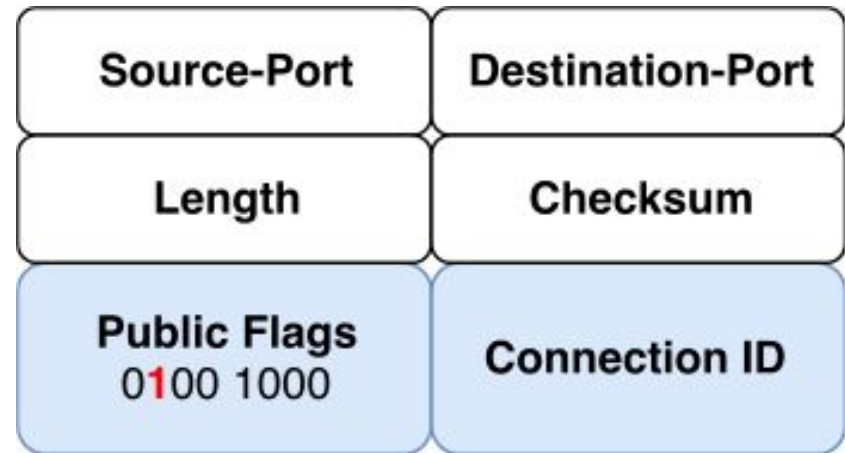
No data ordering based on Packet Numbers!

The Design of MPQUIC

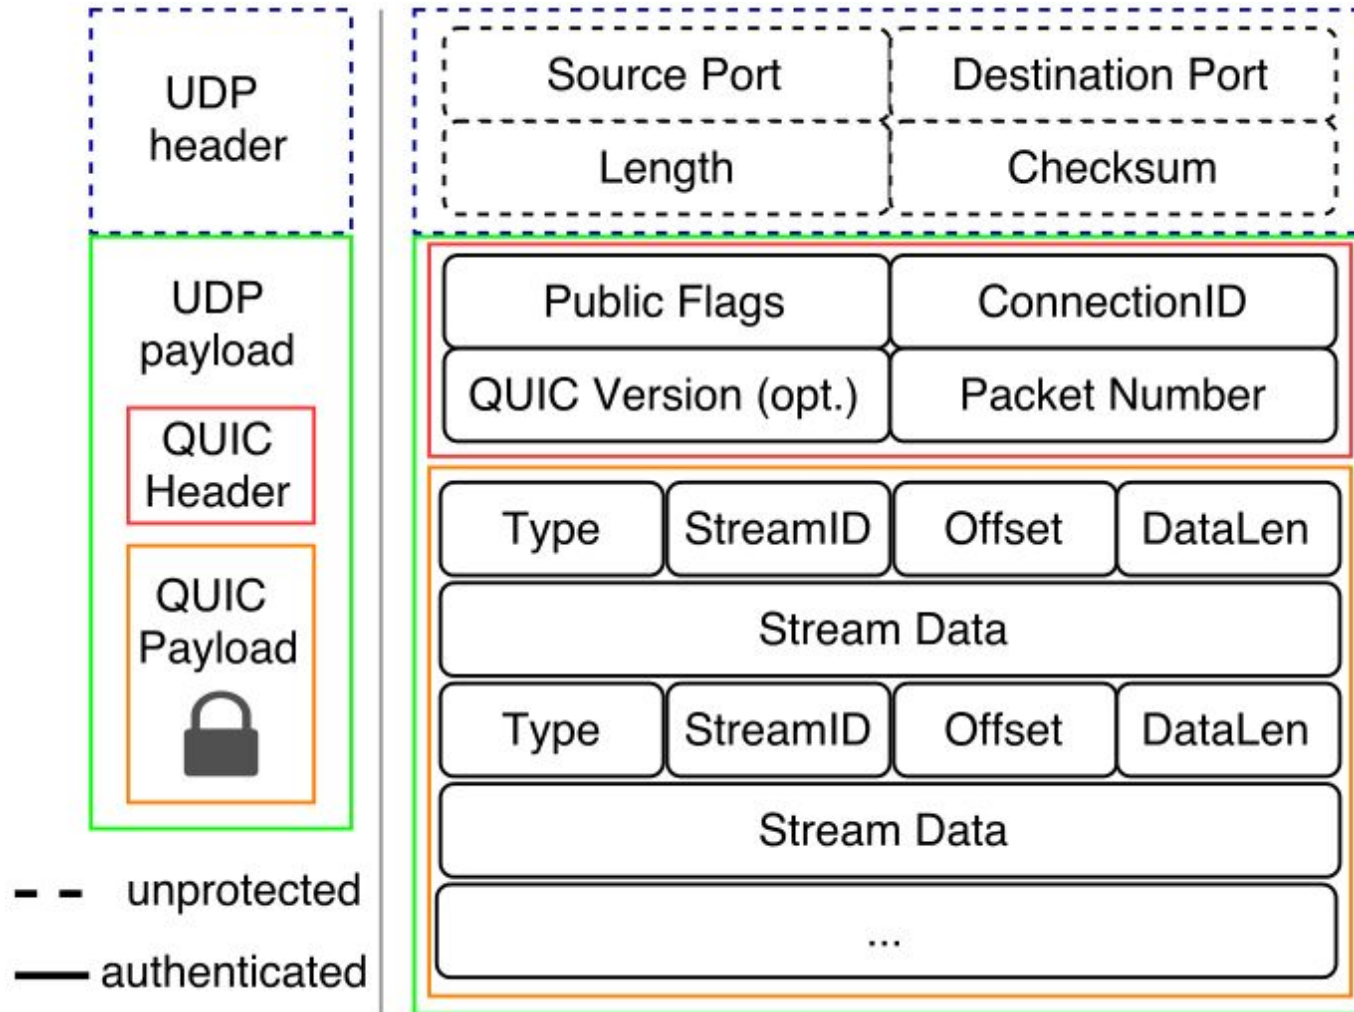
Subflow Establishment: Announcement

We call it Announcement

- Design of an
 - Announcement Packet
 - Announcement-Handshake

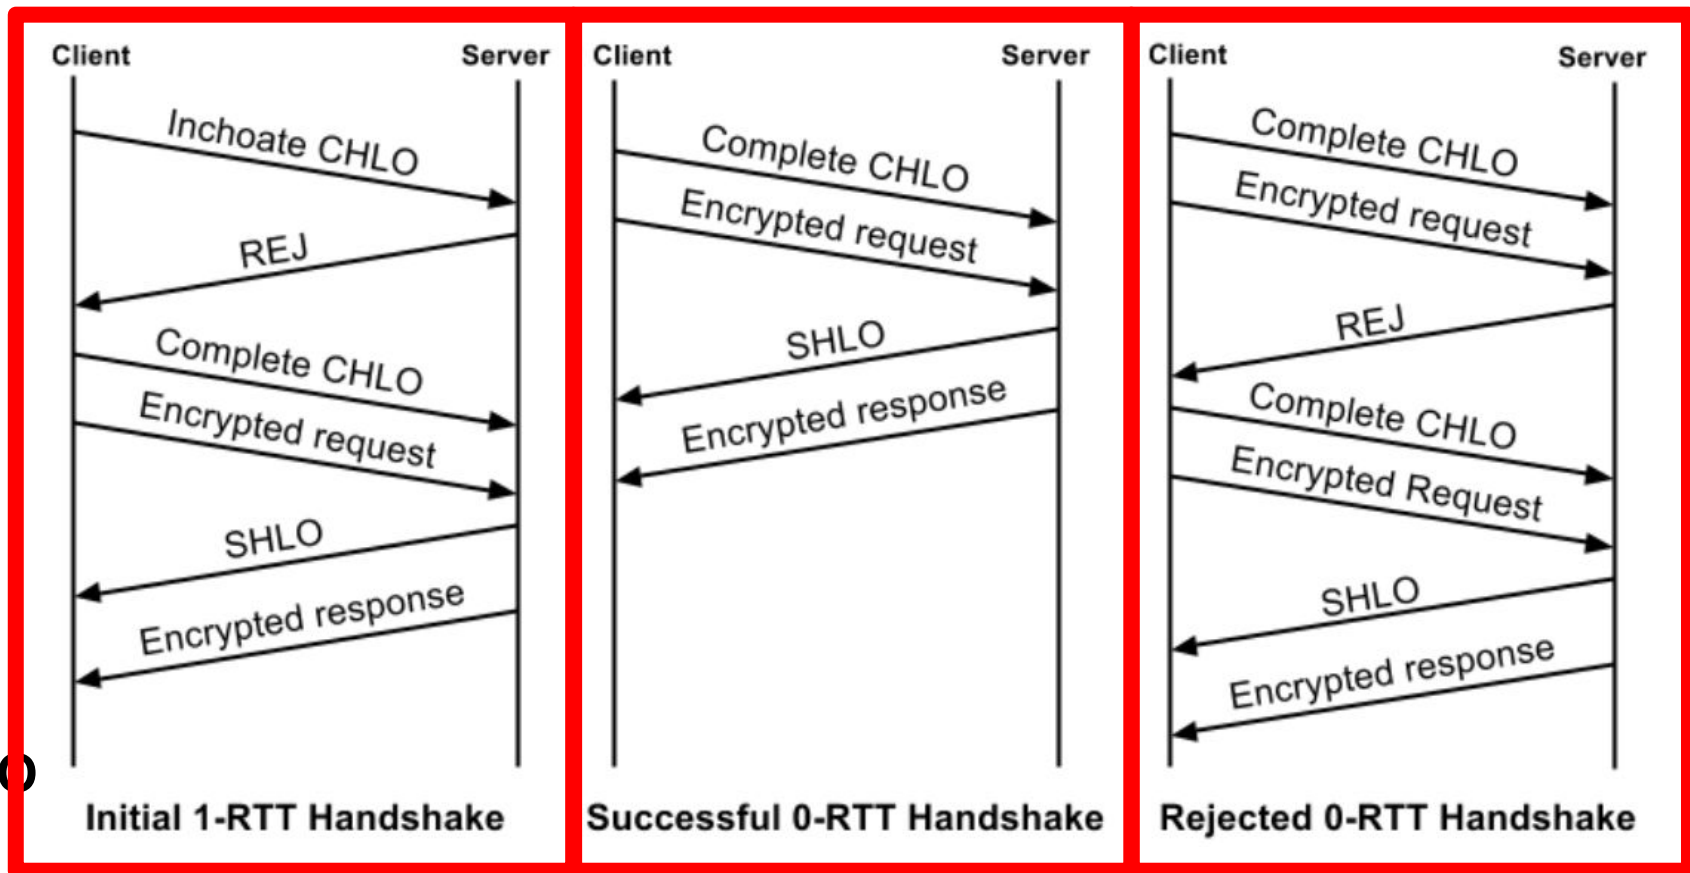


Background: QUIC Packets and Frames



QUIC Connection Establishment


TO



Background: QUIC

QUIC Features:

- Reliable In-Order Delivery
- Congestion and Flow Control
- Fast Connection Establishment
- Stream Multiplexing
- ...





Is QUIC multipath suitable?

Multiplexed Stream

- In-order delivery is not based on Packet Numbers
 - No dependencies between subflows
- Benefits scheduling optimization

Identifying a Connection

- ConnectionID
 - Concept of Connection Migration
 - Fast Connection Establishment
 - Useful for subflows establishment

Deployability

- Shipment with application upgrades
- Inside the UDP Payload
- Authenticated Headers

Reserved Multipath Flag

The Design of MPQUIC

Specifying the individual Components

